

# METU CEng 536

Advanced Unix Programming

## Homework 1

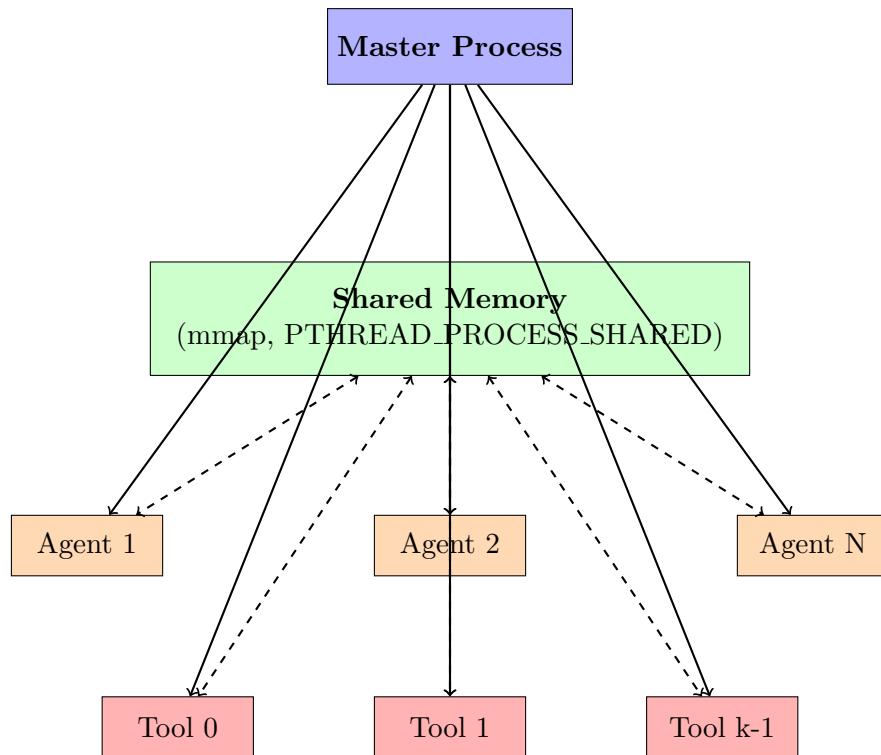
---

# Complete Implementation Guide

---

## IPC Socket-Based Tool Assignment System

Multi-Process Server with Fairness Algorithm



Complete Reference Documentation

Every Line of Code Explained  
From Zero to Full Implementation

Fall 2025  
Version 2.0 - Extended Edition  
November 3, 2025

# **Contents**

## **List of Figures**

## **List of Tables**

## **Listings**



# **Part I**

# **Understanding the Problem**



# Chapter 1

## Assignment Overview

### 1.1 What Are We Building?

You are building a **multi-process server system** that manages gym equipment (tools) with a sophisticated fairness algorithm. This is NOT a simple program - it's a complex distributed system with:

- **Multiple concurrent processes** (master, agents, tools)
- **Shared memory** for inter-process communication
- **Socket-based networking** (both Unix domain and TCP)
- **Complex synchronization** using mutexes and condition variables
- **Fairness algorithm** based on accumulated usage
- **Preemption mechanism** for resource allocation

#### 1.1.1 Real-World Analogy

Imagine a gym with  $k$  treadmills:

1. Customers arrive and want to use treadmills
2. If treadmills are free, customers get assigned immediately
3. If all treadmills are busy, customers wait in a queue
4. The queue is ordered by **fairness**: who has used the treadmill the least total time gets priority
5. There are time limits: after using for  $q$  milliseconds, the system checks if someone with less total usage is waiting
6. After  $Q$  milliseconds, the customer MUST leave the treadmill (unless nobody is waiting)

#### 1.1.2 Technical Complexity Assessment

### 1.2 Project Metrics

#### 1.2.1 Code Statistics

#### 1.2.2 Time Estimates

- **Phase 1 (Foundation):** 8–12 hours
- **Phase 2 (Basic IPC):** 10–15 hours

Component	Difficulty	Why It's Hard
Multi-process architecture	★★★★★	Coordinating multiple processes without deadlocks or race conditions
Shared memory management	★★★★	Cannot use malloc(), must implement custom allocation with integer indices
Process synchronization	★★★★★	PTHREAD_PROCESS_SHARED attributes, condition variable usage
Fairness algorithm	★★★★	Complex decision tree for tool assignment and preemption
Timing simulation	★★★★	Accurate millisecond timing with preemption checks
Socket programming	★★★	Supporting both Unix domain and TCP sockets
Agent threading	★★★★	Two threads per agent: one blocking on socket, one waiting for events

Table 1.1: Technical Difficulty Breakdown

File	Lines	Description
hw1.c	2,500–3,000	Main implementation
heap.c	300	Provided min-heap (modified)
Makefile	20	Build system
<b>Total</b>	<b>3,000</b>	Core implementation
<i>Plus test clients, documentation, etc.</i>		

Table 1.2: Estimated Line Counts

- **Phase 3 (Tool Assignment):** 12–18 hours
- **Phase 4 (Timing & Preemption):** 10–15 hours
- **Phase 5 (Testing & Debug):** 10–15 hours
- **Phase 6 (Documentation):** 5–8 hours

**Total estimated time: 55–83 hours**

### 1.3 Learning Objectives

By completing this assignment, you will master:

1. **Process Management:** fork(), process hierarchy, parent-child relationships
2. **Inter-Process Communication:** Shared memory (mmap), process-shared synchronization
3. **Socket Programming:** Both Unix domain and IPv4 TCP sockets
4. **Multi-threading:** pthread creation, synchronization within and across processes
5. **Synchronization Primitives:** Mutexes, condition variables, avoiding race conditions
6. **Resource Scheduling:** Implementing a fairness algorithm with preemption
7. **Timing and Real-time Constraints:** Millisecond-precision timing
8. **Memory Management:** Custom allocation without malloc() for shared memory
9. **Data Structures:** Heap-based priority queue for fair scheduling

# Chapter 2

# System Architecture Deep Dive

## 2.1 Process Hierarchy

### 2.1.1 The Four Process Types

Our system consists of four distinct process types, each with specific responsibilities:

### 2.1.2 Process Creation Flow

### 2.1.3 Key fork() Concept

fork() creates a complete copy of the calling process:

```
1 // BEFORE fork():
2 // One process exists (parent)
3
4 pid_t pid = fork();
5
6 // AFTER fork():
7 // TWO processes exist!
8
9 if (pid == 0) {
10     // CHILD PROCESS
11     printf("I am the child, PID=%d\n", getpid());
12     printf("My parent PID=%d\n", getppid());
13
14     // Child gets COPY of parent's variables
15     // But shared memory is SHARED (not copied)
16
17 } else if (pid > 0) {
18     // PARENT PROCESS
19     printf("I am the parent, PID=%d\n", getpid());
20     printf("My child PID=%d\n", pid);
21
22 } else {
23     // ERROR
24     perror("fork");
25 }
```

Listing 2.1: fork() Behavior

#### Critical Understanding:

- Normal variables are **copied** (child has separate copy)
- Shared memory (mmap with MAP\_SHARED) is **shared** (both see same memory)
- File descriptors are **copied** (both can use socket)
- Both processes continue from the if statement

<b>Process Type</b>	<b>Count</b>	<b>Responsibilities</b>
Master	1	<ul style="list-style-type: none"> <li>• Create shared memory with mmap()</li> <li>• Set up socket (bind/listen)</li> <li>• Fork <math>k</math> tool processes at startup</li> <li>• Accept client connections in infinite loop</li> <li>• Fork agent process for each client</li> <li>• Never terminates</li> </ul>
Agent	$N$	<ul style="list-style-type: none"> <li>• One per connected client</li> <li>• Allocate customer slot in shared memory</li> <li>• Create two threads: <ul style="list-style-type: none"> <li>– Thread 1: Read from socket (blocking)</li> <li>– Thread 2: Wait for tool events</li> </ul> </li> <li>• Parse commands (REQUEST/REST/REPORT/QUIT)</li> <li>• Update customer state</li> <li>• Send responses to client</li> <li>• Deallocate customer on disconnect</li> <li>• Terminate when client quits</li> </ul>
Tool	$k$	<ul style="list-style-type: none"> <li>• One per tool (e.g., treadmill)</li> <li>• Infinite loop: simulate tool usage timing</li> <li>• Sleep for appropriate duration</li> <li>• Check <math>q</math> and <math>Q</math> limits</li> <li>• Implement preemption logic</li> <li>• Update customer share values</li> <li>• Assign next customer from waiting queue</li> <li>• Send events to agent processes</li> <li>• Never terminates</li> </ul>
Client	$N$	<ul style="list-style-type: none"> <li>• External test programs (not part of hw1)</li> <li>• Connect via socket</li> <li>• Send text commands</li> <li>• Receive and display responses</li> <li>• User-controlled</li> </ul>

Table 2.1: Process Types and Responsibilities

## 2.2 Data Flow Architecture

### 2.2.1 Customer Lifecycle

The complete lifecycle of a customer from connection to termination:

### 2.2.2 REQUEST Command Detailed Flow

When a customer sends REQUEST 5000:

1. **Agent socket thread** receives the command
2. Thread calls `handle_request(customer_idx, 5000)`
3. `handle_request()` locks `global_mutex`
4. Check customer's current state:
  - If RESTING: Update `resting_customers` count
  - If already USING: Error (invalid state)
5. Set `request_duration = 5000`, `remaining_duration = 5000`
6. **Tool Assignment Logic:**
  - Search for free tool: `find_free_tool()`
  - If found: Immediately assign with `assign_tool_to_customer()`
  - If not found:
    - Check if can preempt: `find_preemption_candidate()`
    - If can preempt: Remove current user, assign to new customer
    - If cannot preempt: Add to waiting queue (`heap_insert()`)
7. Unlock `global_mutex`
8. If assigned: Agent notify thread wakes up and sends message to client
9. If in queue: Agent waits for tool to become available

## 2.3 Synchronization Strategy

### 2.3.1 Why Synchronization is Essential

Without proper synchronization, consider this scenario:

```

1 // Shared: int counter = 0;
2
3 // Process A (Tool 0):           Process B (Tool 1):
4 tmp = counter; // Read: 0       tmp = counter; // Read: 0
5 tmp = tmp + 1; // Compute: 1    tmp = tmp + 1; // Compute: 1
6 counter = tmp; // Write: 1     counter = tmp; // Write: 1
7
8 // RESULT: counter = 1 (WRONG! Should be 2)

```

Listing 2.2: Race Condition Example - NO MUTEX

With mutex:

```
1 // Process A:                                Process B:  
2 pthread_mutex_lock(&mutex);                // BLOCKED - waiting...  
3 tmp = counter;    // Read: 0                 // Still waiting...  
4 tmp = tmp + 1;   // Compute: 1              // Still waiting...  
5 counter = tmp;   // Write: 1                // Still waiting...  
6 pthread_mutex_unlock(&mutex);  
7  
8  
9  
10  
11  
12 // RESULT: counter = 2 (CORRECT!)
```

Listing 2.3: Correct Synchronization - WITH MUTEX

### 2.3.2 Our Synchronization Primitives

**Critical Rule:** *Always lock global\_mutex before accessing ANY shared memory data!*

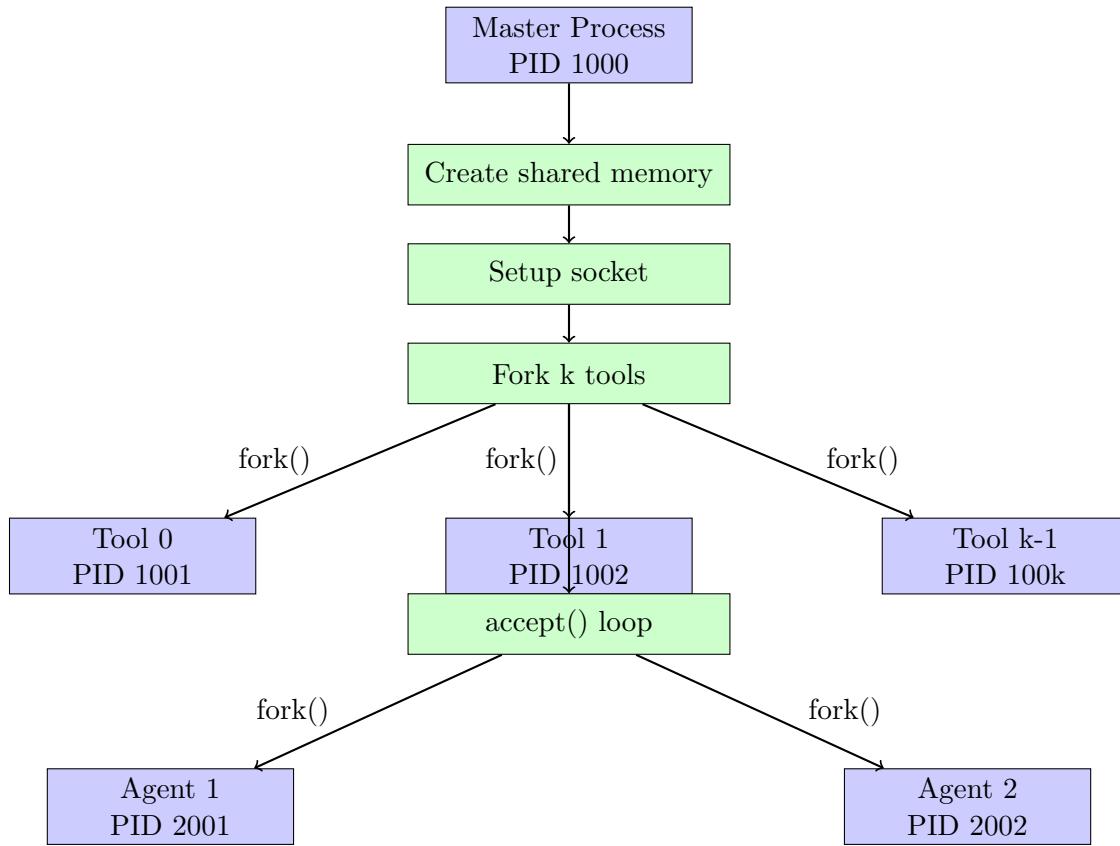


Figure 2.1: Process Creation Sequence

Primitive	Count	Purpose
global_mutex	1	Protects ALL shared memory access
new_customer_cond	1	Signals tool processes when new customer arrives
agent_cond	$N$	One per customer, tool notifies agent of events
tool_cond	$k$	One per tool, agent wakes tool when customer assigned

Table 2.2: Synchronization Primitives

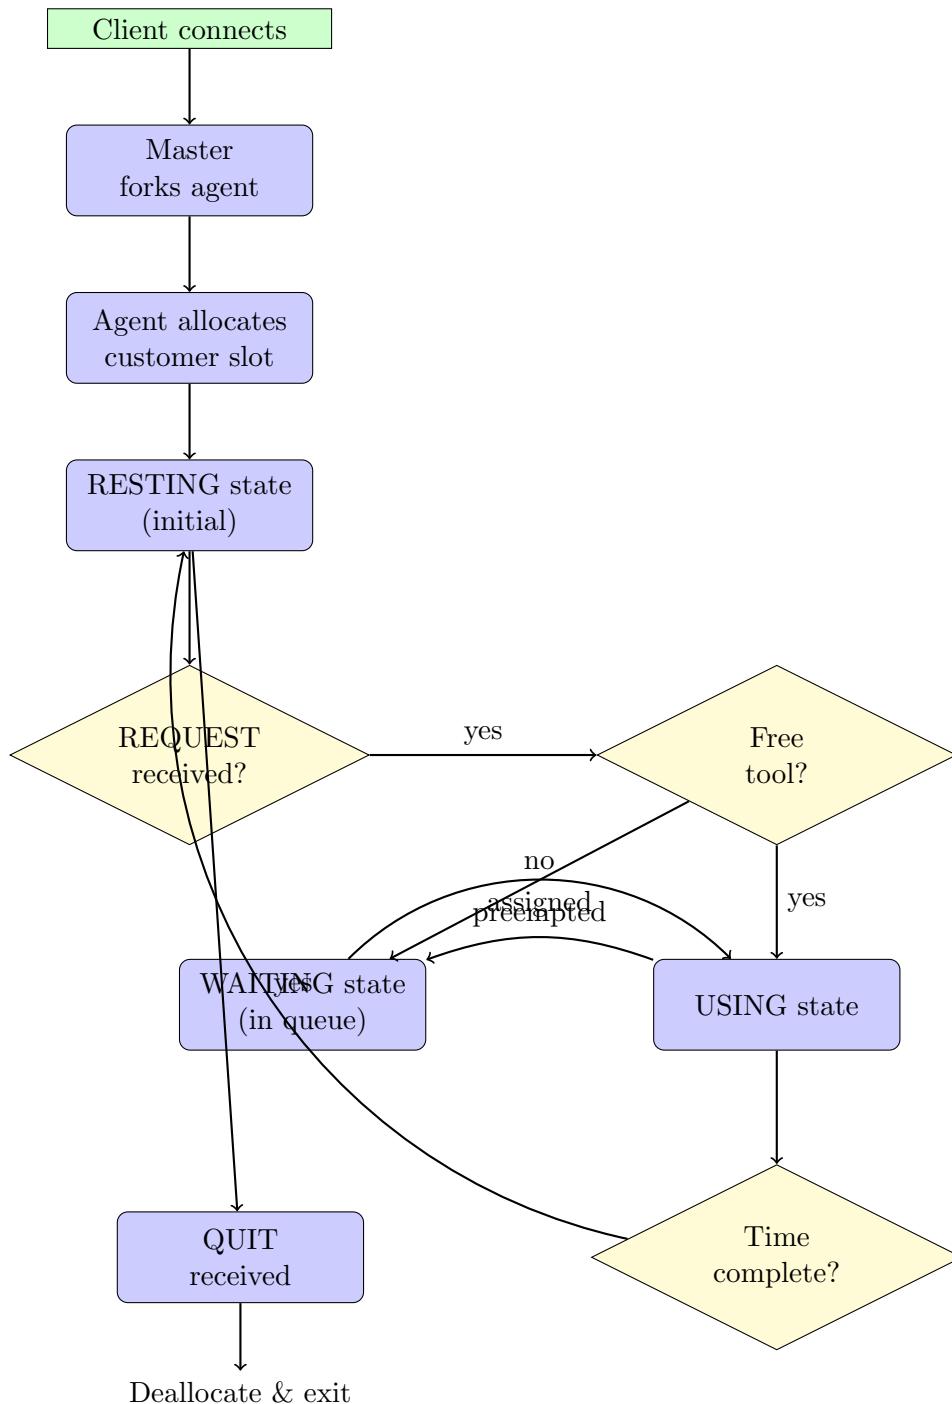


Figure 2.2: Complete Customer State Machine

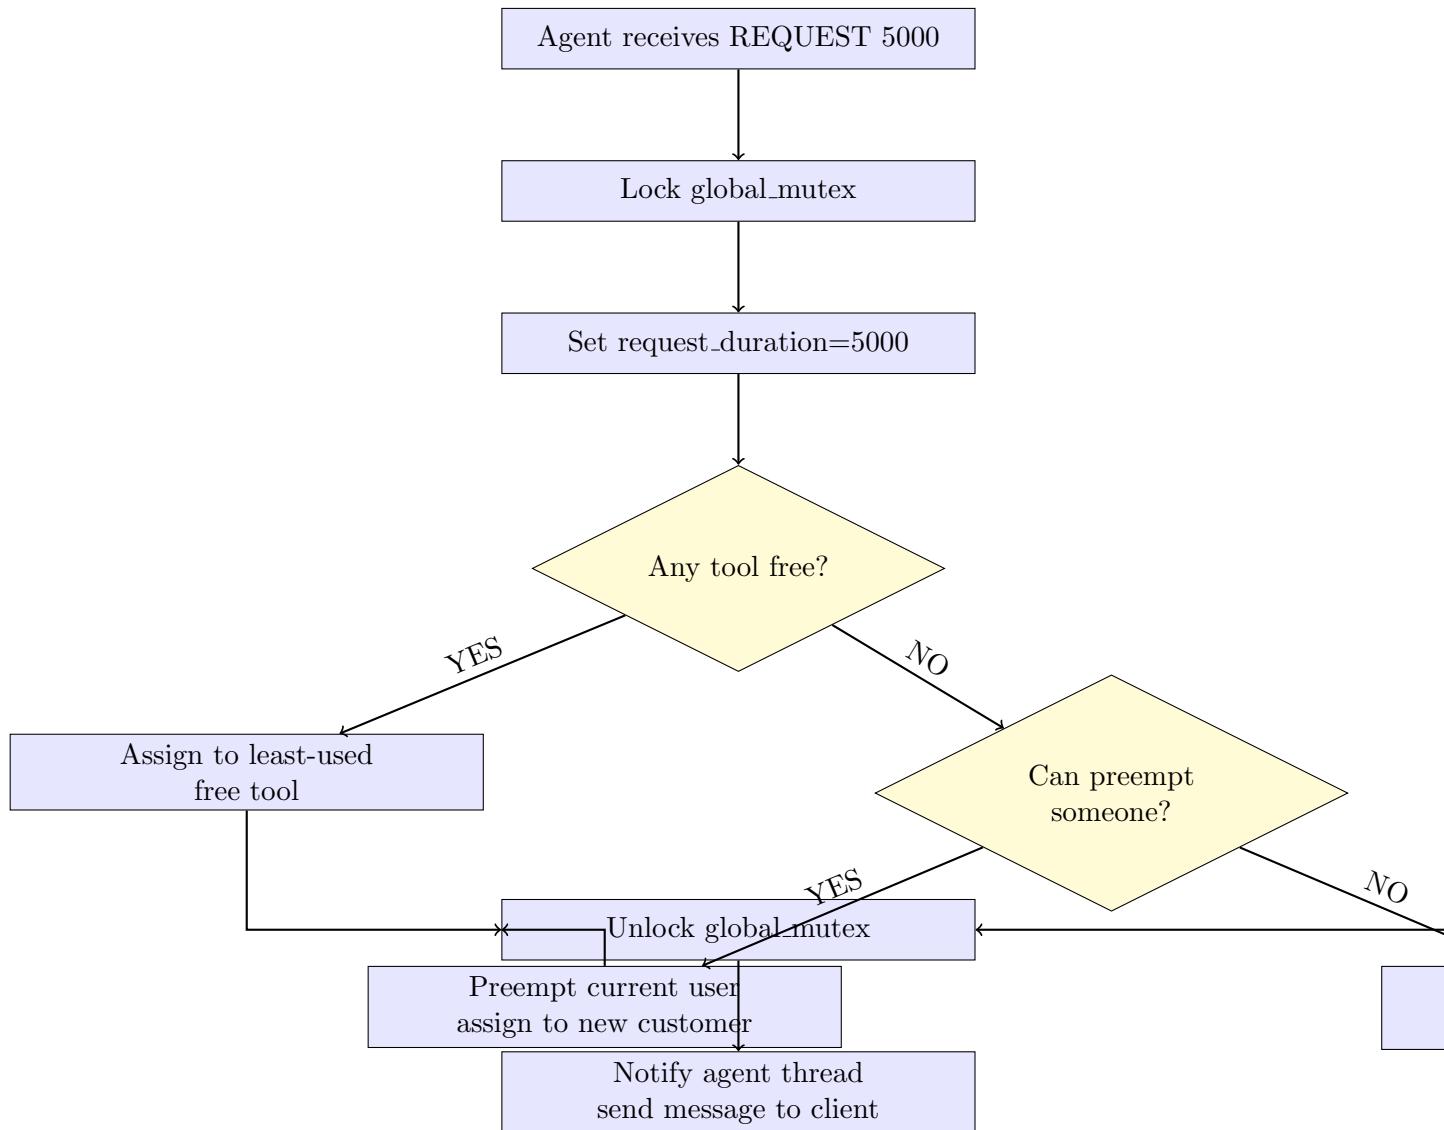


Figure 2.3: REQUEST Command Processing Flow



## **Part II**

# **Data Structures in Detail**



# Chapter 3

## Core Structures

### 3.1 Customer Structure - Line by Line

The Customer structure is the heart of the fairness algorithm. Let's examine every field:

```
1 typedef struct {
2     //=====
3     // IDENTIFICATION
4     //=====
5
6     int customer_id;           // Agent process PID
7                         // Why: Uniquely identify each customer
8                         // Used in: printf output, debugging
9
10    int is_allocated;         // 1 = slot in use, 0 = free
11                         // Why: Track which slots are occupied
12                         // Used in: allocation/deallocation
13
14    //=====
15    // STATE MANAGEMENT
16    //=====
17
18    CustomerState state;      // Current state enum
19                         // Why: Track customer's current activity
20                         // Values: RESTING / WAITING / USING / DELETED
21                         // Used in: state machine transitions
22
23    //=====
24    // FAIRNESS TRACKING (MOST CRITICAL!)
25    //=====
26
27    double share;             // Total accumulated usage time (milliseconds)
28                         // Why: THIS IS THE FAIRNESS METRIC
29                         // Lower share = higher priority in queue
30                         // Used in: heap sorting, preemption decisions
31                         // Example: Customer with share=5000 has used
32                         //           tools for 5 seconds total
33
34    //=====
35    // CURRENT REQUEST
36    //=====
37
38    int request_duration;     // Duration requested in REQUEST command
39                         // Why: Track how much time customer wants
40                         // Used in: timing calculations
41                         // Example: REQUEST 5000 sets this to 5000
42
43    int remaining_duration;   // How much time left to use
44                         // Why: Track progress through request
45                         // Used in: session completion check
```

```

46                                     // Example: After 2s, this becomes 3000
47
48 //=====
49 // TOOL USAGE
50 //=====
51
52 int current_tool;           // Which tool is customer using?
53                                         // Why: Track tool assignment
54                                         // Value: 0 to k-1, or -1 if not using
55                                         // Used in: tool release, reporting
56
57 long long session_start;    // Timestamp when tool usage began (ms)
58                                         // Why: Calculate elapsed time
59                                         // Used in: timing, duration calculations
60                                         // Example: start=1000, now=3000      elapsed=2000ms
61
62 //=====
63 // WAITING QUEUE
64 //=====
65
66 long long wait_start;       // Timestamp when entered waiting queue (ms)
67                                         // Why: Calculate how long customer has waited
68                                         // Used in: REPORT command duration column
69
70 int heap_index;             // Index in GLB array (from heap.c)
71                                         // Why: Link customer to heap node
72                                         // Used in: heap operations
73                                         // Value: Same as customer array index
74
75 //=====
76 // AGENT COMMUNICATION
77 //=====
78
79 pthread_cond_t agent_cond; // Condition variable for events
80                                         // Why: Tool wakes up agent thread
81                                         // Used in: pthread_cond_wait/signal
82
83 int event_pending;          // Flag: new event available? (1=yes, 0=no)
84                                         // Why: Prevent spurious wakeups
85                                         // Used in: while(!event_pending) wait()
86
87 EventType event_type;        // What kind of event occurred?
88                                         // Why: Tell agent what happened
89                                         // Values: ASSIGNED / REMOVED / COMPLETED
90
91 int event_tool_id;          // Which tool is this event about?
92                                         // Why: Include tool number in message
93                                         // Used in: printf output
94
95 } Customer;

```

Listing 3.1: Customer Structure - Complete

### 3.1.1 Share Field Deep Dive

The share field is the most important field for fairness.

**Calculation Examples:**

**Priority Ordering:**

Highest Priority	Customer A share=300	Customer B share=550	Customer C share=800	Lowest Priority
------------------	----------------------	----------------------	----------------------	-----------------

**In Waiting Queue:** Customer A will be assigned tool first.

## 3.2 Tool Structure Explained

```

1 typedef struct {
2     //=====
3     // IDENTIFICATION
4     //=====
5
6     int tool_id;           // Tool number (0 to k-1)
7                         // Why: Identify this tool
8                         // Used in: printf output, reports
9
10    //=====
11    // USAGE STATISTICS
12    //=====
13
14    long long total_usage;   // Total milliseconds this tool has been used
15                         // Why: Balance load across tools
16                         // Used in: find_free_tool() selection
17                         // Example: Tool 0 used for 50s, Tool 1 for 45s
18                         // Assign to Tool 1 (less used)
19
20    //=====
21    // CURRENT SESSION
22    //=====
23
24    int current_user;      // Customer index (-1 if free)
25                         // Why: Track who is using this tool
26                         // Value: Index in customers[] array, or -1
27
28    int current_usage;     // Milliseconds in THIS session
29                         // Why: Check q and Q limits
30                         // Used in: preemption logic
31                         // Example: If current_usage >= Q, must preempt
32
33    long long session_start; // When did current session start? (ms)
34                         // Why: Calculate elapsed time
35                         // Used in: timing calculations
36
37    //=====
38    // SYNCHRONIZATION
39    //=====
40
41    pthread_cond_t tool_cond; // Condition variable for this tool
42                         // Why: Wake up tool thread
43                         // Used in: When customer assigned or freed
44
45    int tool_should_exit;   // Flag for cleanup (1=exit, 0=run)
46                         // Why: Graceful shutdown
47                         // Used in: Signal handling
48
49 } ToolInfo;

```

Listing 3.2: Tool Structure - Complete

### 3.2.1 Tool Selection Algorithm

When multiple tools are free, which one do we assign?

**Rule:** Select the tool with the **lowest total\_usage**. If tied, select **lowest tool.id**.

**Example:**

**Decision Process:**

1. Tools 0, 1, 3 are free
2. Tools 1 and 3 have lowest usage (45,000 ms)

3. Tool 1 has lower ID than Tool 3

#### 4. Result: Assign to Tool 1

```

1 int find_free_tool(void) {
2     int best_tool = -1;
3     long long min_usage = LLONG_MAX;
4
5     for (int i = 0; i < shm->num_tools; i++) {
6         if (shm->tools[i].current_user == -1) {
7             // Tool is free
8             if (shm->tools[i].total_usage < min_usage) {
9                 min_usage = shm->tools[i].total_usage;
10                best_tool = i;
11            } else if (shm->tools[i].total_usage == min_usage) {
12                // Tie: select lower ID
13                if (i < best_tool || best_tool == -1) {
14                    best_tool = i;
15                }
16            }
17        }
18    }
19
20    return best_tool; // -1 if no free tools
21 }
```

Listing 3.3: Tool Selection Code Snippet

### 3.3 Shared Memory Structure

```

1 typedef struct {
2     //=====
3     // CUSTOMER POOL (Memory Management)
4     //=====
5
6     Customer customers[MAX_CUSTOMERS];
7     // Why: Fixed-size array of all customer slots
8     // Size: 1,000,000 slots
9     // Note: Cannot use malloc() in shared memory!
10
11    int free_customer_slots[MAX_CUSTOMERS];
12    // Why: Free list for available slots
13    // Contains: Indices of free slots
14    // Example: [0, 1, 2, ..., 999999] initially
15    // After allocating slot 5: [0,1,2,3,4,6,7,...]
16
17    int free_customer_count;
18    // Why: How many free slots remain?
19    // Initial value: 1,000,000
20    // After 3 customers: 999,997
21
22    //=====
23    // WAITING QUEUE
24    //=====
25
26    int waiting_count;
27    // Why: Track number of customers in waiting queue
28    // Note: Actual heap is in heap.c (GLB, HEAP_ARRAY)
29    // Used in: Reports, queue empty checks
30
31    //=====
32    // TOOL INFORMATION
33    //=====
```

```

34     ToolInfo tools[MAX_TOOLS];
35     // Why: Fixed-size array of tool structures
36     // Size: 100 slots (max k=100)
37
38     int num_tools;
39     // Why: Actual number of tools (k from command line)
40     // Example: ./hw1 ... 3 means num_tools=3
41
42     //=====
43     // GLOBAL STATISTICS
44     //=====
45
46     int total_customers;
47     // Why: Count of active customers (allocated slots)
48     // Used in: Average share calculation, reports
49
50     int resting_customers;
51     // Why: Count in RESTING state
52     // Used in: Reports
53
54     double total_share;
55     // Why: Sum of all customer shares
56     // Used in: Calculate average for new customers
57     // Example: 3 customers with shares 5000,8000,6000
58     //           total_share = 19000
59     //           New customer gets 19000/3 = 6333.33
60
61     //=====
62     // SYNCHRONIZATION
63     //=====
64
65     pthread_mutex_t global_mutex;
66     // Why: THE main mutex protecting all shared data
67     // Rule: ALWAYS lock before accessing shared memory!
68     // Type: PTHREAD_PROCESS_SHARED
69
70     pthread_cond_t new_customer_cond;
71     // Why: Signal all tools when customer enters queue
72     // Used in: Wake idle tools
73
74     //=====
75     // CONFIGURATION
76     //=====
77
78     int q;
79     // Why: Minimum uninterrupted time (milliseconds)
80     // From: Command line argument
81     // Used in: Check preemption after q ms
82
83     int Q;
84     // Why: Maximum uninterrupted time (milliseconds)
85     // From: Command line argument
86     // Used in: Force preemption after Q ms
87
88     long long start_time;
89     // Why: Server start timestamp
90     // Used in: Calculate uptime, debugging
91
92 } SharedMemory;

```

Listing 3.4: SharedMemory Structure - Complete

### 3.3.1 Memory Management Strategy

**Problem:** Cannot use `malloc()` in shared memory!

**Solution:** Pre-allocate fixed-size arrays and use free list.

**Allocation Algorithm:**

```
1 // Get free slot
2 int idx = free_customer_slots[--free_customer_count];
3
4 // Use slot
5 customers[idx] = /* initialize */;
```

**Deallocation Algorithm:**

```
1 // Return slot to free list
2 free_customer_slots[free_customer_count++] = idx;
```

Event	Elapsed (ms)	New Share	Calculation
Customer arrives	—	0	First customer
Uses tool for 3s	3000	3000	$0 + 3000$
Uses tool for 2s more	2000	5000	$3000 + 2000$
Rests for 10s	0	5000	No change (not using tool)
Uses tool for 1s	1000	6000	$5000 + 1000$

Table 3.1: Share Calculation Examples

Tool ID	Total Usage (ms)	Current User	Selection
0	50,000	-1 (FREE)	
1	45,000	-1 (FREE)	← SELECT THIS!
2	52,000	Customer 5	Not available
3	45,000	-1 (FREE)	

Table 3.2: Tool Selection Example

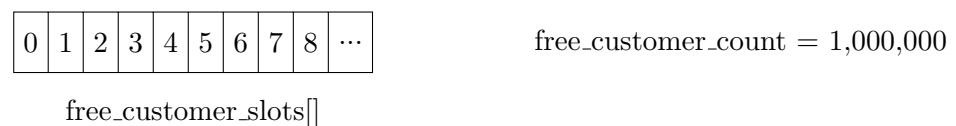
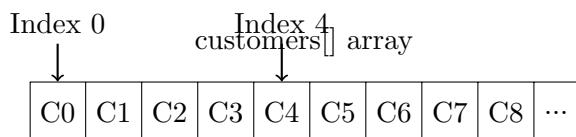


Figure 3.1: Customer Pool Memory Layout



## **Part III**

# **Complete Implementation - Every Function**



# Chapter 4

## Phase 1: Foundation

### 4.1 Project Structure Setup

#### 4.1.1 Directory Organization

Create this exact structure:

```
~/ceng536-hw1/
hw1/                      # ← Submission directory
    hw1.c                  # ← Your main code
    heap.c                 # ← Instructor's code (modified)
    Makefile                # ← Build system

test/                      # ← Testing tools (DON'T submit)
    test_client.py          # ← Python test client
    test_client.c           # ← C test client (optional)
    test_scenarios.sh       # ← Automated tests

docs/                      # ← Your notes
    notes.md
    algorithm.txt

original/                  # ← Backup
    heap.c                 # ← Original from instructor
    rb.c                   # ← (not using this)
    ceng536-hw1.pdf

submission/                # ← Final tar.gz goes here
    yourname_id.tar.gz
```

#### 4.1.2 Step-by-Step Setup

Commands to run:

```
1 # 1. Create directory structure
2 cd ~
3 mkdir -p ceng536-hw1/{hw1,test,docs,original,submission}
4
5 # 2. Copy instructor's files
6 cd ~/ceng536-hw1
7 cp ~/Downloads/heap.c original/
8 cp ~/Downloads/rb.c original/
9 cp ~/Downloads/ceng536-hw1.pdf original/
10
11 # 3. Copy heap.c to working directory
```

```

12 cp original/heap.c hw1/
13
14 # 4. Verify structure
15 tree ceng536-hw1

```

### 4.1.3 Modify heap.c

Open `hw1/heap.c` and make these changes:

#### Change 1: MAX\_NODES

Find line 32:

```
1 #define MAX_NODES 1024
```

Change to:

```
1 #define MAX_NODES 1000000
```

#### Change 2: Remove main()

Find line 509 (the main function):

```

1 int main() {
2     // ... test code ...
3     return 0;
4 }
```

Comment it out:

```

1 /*
2 int main() {
3     // ... test code ...
4     return 0;
5 }
6 */
```

**Why remove main()?** Because your `hw1.c` will have its own `main()`!

### 4.1.4 Create Makefile

Create `hw1/Makefile`:

```

1 # Makefile for CEng 536 HW1
2
3 CC = gcc
4 CFLAGS = -Wall -Wextra -Werror -g -pthread -std=c11 \
           -D_POSIX_C_SOURCE=200809L
5
6 LDFLAGS = -lpthread -lrt
7
8 TARGET = hw1
9 SRCS = hw1.c heap.c
10 OBJS = $(SRCS:.c=.o)
11
12 # Main target
13 $(TARGET): $(OBJS)
14     $(CC) $(CFLAGS) -o $(TARGET) $(OBJS) $(LDFLAGS)
15
16 # Compile .c to .o
17 %.o: %.c
18     $(CC) $(CFLAGS) -c $< -o $@
19
20 # Clean
21 clean:
22     rm -f $(TARGET) $(OBJS)
23
24 # Debug build
25 debug: CFLAGS += -DDEBUG -O0

```

```

26 debug: clean $(TARGET)
27
28 # Run with example parameters
29 run-test: $(TARGET)
30     ./$(TARGET) @/tmp/gym.sock 1000 5000 3
31
32 # Check for memory leaks
33 valgrind: $(TARGET)
34     valgrind --leak-check=full --show-leak-kinds=all \
35         ./$(TARGET) @/tmp/gym.sock 1000 5000 2
36
37 .PHONY: clean debug run-test valgrind

```

Test the Makefile:

```

1 cd ~/ceng536-hw1/hw1
2 make clean
3 # Should output: rm -f hw1 *.o

```

## 4.2 hw1.c: Basic Skeleton

Create hw1/hw1.c with this initial structure:

```

1 /*=====
2 * CEng 536 - Advanced Unix Programming
3 * Homework 1: IPC Gym Tool Assignment System
4 *
5 * Multi-process server with fairness algorithm
6 =====*/
7
8 #include <stdio.h>
9 #include <stdlib.h>
10 #include <string.h>
11 #include <unistd.h>
12 #include <errno.h>
13 #include <signal.h>
14 #include <ctype.h>
15
16 // Socket headers
17 #include <sys/socket.h>
18 #include <sys/un.h>           // Unix domain sockets
19 #include <netinet/in.h>       // Internet sockets
20 #include <arpa/inet.h>
21
22 // Process & threading
23 #include <sys/types.h>
24 #include <sys/wait.h>
25 #include <pthread.h>
26
27 // Shared memory
28 #include <sys/mman.h>
29
30 // Timing
31 #include <time.h>
32 #include <sys/time.h>
33
34 // Limits
35 #include <limits.h>
36 #include <float.h>
37
38 /*=====
39 * HEAP INTERFACE (from heap.c)
40 =====*/
41
42 // Heap functions

```

```
43 extern int heap_insert(int nodeindex);
44 extern int heap_pop(void);
45 extern int heap_delete(int nodeindex);
46 extern int heap_size;
47
48 // Heap node structure
49 struct Node {
50     double key;
51     int heap_index;
52 };
53 extern struct Node GLB[];
54 extern int HEAP_ARRAY[];
55
56 /*=====
57 * CONSTANTS
58 =====*/
59
60 #define MAX_CUSTOMERS 1000000
61 #define MAX_TOOLS 100
62 #define BUFFER_SIZE 4096
63 #define NIL -1
64
65 /*=====
66 * DEBUG LOGGING
67 =====*/
68
69 #ifdef DEBUG
70 #define LOG(fmt, ...) \
71     do { \
72         struct timespec ts; \
73         clock_gettime(CLOCK_MONOTONIC, &ts); \
74         fprintf(stderr, "[%ld.%03ld][PID:%d] " fmt "\n", \
75                 ts.tv_sec, ts.tv_nsec/1000000, getpid(), \
76                 ##__VA_ARGS__); \
77         fflush(stderr); \
78     } while(0)
79 #else
80 #define LOG(fmt, ...) do {} while(0)
81 #endif
82
83 /*=====
84 * ENUMERATIONS
85 =====*/
86
87 typedef enum {
88     CUSTOMER_STATE_RESTING,
89     CUSTOMER_STATE_WAITING,
90     CUSTOMER_STATE_USING,
91     CUSTOMER_STATE_DELETED
92 } CustomerState;
93
94 typedef enum {
95     EVENT_NONE = 0,
96     EVENT_TOOL_ASSIGNED,
97     EVENT_TOOL_REMOVED,
98     EVENT_TOOL_COMPLETED
99 } EventType;
100
101 /*=====
102 * STRUCTURE DECLARATIONS
103 =====*/
104
105 typedef struct Customer Customer;
106 typedef struct ToolInfo ToolInfo;
107 typedef struct SharedMemory SharedMemory;
108 typedef struct ServerAddress ServerAddress;
```

```
109 typedef struct AgentContext AgentContext;
110
111 /*=====
112 * CUSTOMER STRUCTURE
113 =====*/
114
115 struct Customer {
116     // Identification
117     int customer_id;
118     int is_allocated;
119
120     // State
121     CustomerState state;
122
123     // Fairness
124     double share;
125
126     // Request
127     int request_duration;
128     int remaining_duration;
129
130     // Tool usage
131     int current_tool;
132     long long session_start;
133
134     // Waiting queue
135     long long wait_start;
136     int heap_index;
137
138     // Communication
139     pthread_cond_t agent_cond;
140     int event_pending;
141     EventType event_type;
142     int event_tool_id;
143 };
144
145 /*=====
146 * TOOL STRUCTURE
147 =====*/
148
149 struct ToolInfo {
150     int tool_id;
151
152     long long total_usage;
153
154     int current_user;
155     int current_usage;
156     long long session_start;
157
158     pthread_cond_t tool_cond;
159     int tool_should_exit;
160 };
161
162 /*=====
163 * SHARED MEMORY STRUCTURE
164 =====*/
165
166 struct SharedMemory {
167     Customer customers[MAX_CUSTOMERS];
168     int free_customer_slots[MAX_CUSTOMERS];
169     int free_customer_count;
170
171     int waiting_count;
172
173     ToolInfo tools[MAX_TOOLS];
174     int num_tools;
```

```
175     int total_customers;
176     int resting_customers;
177     double total_share;
178
179     pthread_mutex_t global_mutex;
180     pthread_cond_t new_customer_cond;
181
182     int q;
183     int Q;
184     long long start_time;
185 };
186
187
188 /*=====
189 * SERVER ADDRESS
190 =====*/
191
192 struct ServerAddress {
193     int is_unix;
194     char path[256];
195     char ip[64];
196     int port;
197 };
198
199 /*=====
200 * AGENT CONTEXT
201 =====*/
202
203 struct AgentContext {
204     int socket_fd;
205     int customer_idx;
206     int should_exit;
207 };
208
209 /*=====
210 * GLOBAL VARIABLES
211 =====*/
212
213 static SharedMemory *shm = NULL;
214
215 /*=====
216 * FUNCTION PROTOTYPES
217 =====*/
218
219 // Utility
220 long long get_current_time_ms(void);
221 ServerAddress parse_address(const char *addr_str);
222
223 // Shared memory
224 void setup_shared_memory(int q, int Q, int k);
225 void setup_heap(void);
226
227 // Customer management
228 int allocate_customer(int customer_id);
229 void deallocate_customer(int customer_idx);
230
231 // Tool assignment
232 int find_free_tool(void);
233 int find_preemption_candidate(double new_customer_share);
234 void assign_tool_to_customer(int customer_idx, int tool_id);
235 void assign_next_from_queue(int tool_id);
236
237 // Request handlers
238 void handle_request(int customer_idx, int duration);
239 void handle_rest(int customer_idx);
240 void handle_report(int socket_fd);
```

```
241 // Tool process
242 void* tool_thread_func(void* arg);
243 void handle_session_complete(int tool_id);
244 void handle_Q_limit(int tool_id);
245 void check_preemption(int tool_id);
246
247 // Agent process
248 void agent_process(int client_socket);
249 void* agent_socket_thread(void* arg);
250 void* agent_notify_thread(void* arg);
251
252 // Socket
253 int create_server_socket(ServerAddress addr);
254
255 // Main
256 int main(int argc, char *argv[]);
257
258 /*=====
259  * IMPLEMENTATION BEGINS HERE
260  * =====*/
261
```

Listing 4.1: hw1.c - Initial Skeleton (Part 1/10)

Now test compilation:

```
1 cd ~/ceng536-hw1/hw1
2 make clean
3 make hw1
4 # Expected: errors because functions not implemented yet
```

This is normal! We'll implement each function step by step.